

# ECG Heartbeat Classification

In this practice we will detect cases of cardiovascular disease through the analysis of heartbeats

This exercise is based on [paper](#) that solves the problem we are facing.

## ▼ 1. Data Analysis

### ▼ Library Import

```
1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import numpy as np
```

### ▼ Dataset Import

```
1 # Connect with Google Drive
2 from google.colab import drive
3 drive.mount('/content/drive')
```

Mounted at /content/drive

```
1 data = pd.read_csv('/content/drive/MyDrive/datasets/DL1_ECD/mitbih_train.csv', header=None)
```

```
1 data.shape
```

(87554, 188)

```
1 data.head()
```

Each column represents an elctrocardiogram reading (at 125hz). In total ther are 187 readings, in this columns we have about a second and half off keystrokes. The last column contains the category to which these keystrokes belong. In total there are five, each represented by a number:


- Normal:0
- Premature arrhythmia:1
- Ventricular premature contraction or Ventricular escape:2
- Fusion of ventricular and normal contraction:3
- Resuscitation, fusion of normal and resucitation or unclassifiable:4

## ▼ Data Distribution

```
1 data.describe()
```


	0	1	2	3	4	
<b>count</b>	87554.000000	87554.000000	87554.000000	87554.000000	87554.000000	87554.0
<b>mean</b>	0.890360	0.758160	0.423972	0.219104	0.201127	0.2
<b>std</b>	0.240909	0.221813	0.227305	0.206878	0.177058	0.1
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.0
<b>25%</b>	0.921922	0.682486	0.250969	0.048458	0.082329	0.0
<b>50%</b>	0.991342	0.826013	0.429472	0.166000	0.147878	0.1
<b>75%</b>	1.000000	0.910506	0.578767	0.341727	0.258993	0.2
<b>max</b>	1.000000	1.000000	1.000000	1.000000	1.000000	1.0

8 rows × 188 columns



```
1 plt.plot(data.iloc[3])
```

```
[<matplotlib.lines.Line2D at 0x7f59e43072d0>]
```



## 2. Data Pocessing

Now that we have visualized our data, let's work with it. First, to divide them into input and output.

```
1 # Converting the Dataset to a Numpy array
2 M = data.values
3 X = M[:, :-1] # Matriz M without the last column
4 y = M[:, -1].astype(int) # The last column of M
5 y
```

```
array([0, 0, 0, ..., 4, 4, 4])
```

Arrays are created with the indices of the examples that belong to each category. can be used [np.argwhere](#) and [np.flatten](#)

```
1 C0 = np.argwhere(y == 0).flatten()
2 C1 = np.argwhere(y == 1).flatten()
3 C2 = np.argwhere(y == 2).flatten()
4 C3 = np.argwhere(y == 3).flatten()
5 C4 = np.argwhere(y == 4).flatten()
6 C4
```

```
array([81123, 81124, 81125, ..., 87551, 87552, 87553])
```

```
1 # count how many examples we have of each category
2 u = {'N': C0, 'S': C1, 'V': C2, 'F': C3, 'Q': C4}
3 www = []
4 for k in u:
5     print('Hay {} muestras de la categoría {}'.format(len(u[k]), k))
6     www.append(len(u[k]))
```

```
Hay 72471 muestras de la categoría N
Hay 2223 muestras de la categoría S
Hay 5788 muestras de la categoría V
Hay 641 muestras de la categoría F
Hay 6431 muestras de la categoría Q
```


```
1 www
```

```
[72471, 2223, 5788, 641, 6431]
```

To better see how many of each type we have, we are going to make a bar graph and a plot

```
1 labels = list(u.keys())  
2 plt.bar(labels, www)
```

<BarContainer object of 5 artists>



We can print an electrocardiogram of each type

```
1 plt.figure(figsize=(15,8))  
2 plt.plot(M[C0[20]], label='N')  
3 plt.plot(M[C1[20]], label='S')  
4 plt.plot(M[C2[20]], label='V')  
5 plt.legend()  
6 plt.title('Different classes')
```

```
Text(0.5, 1.0, 'Different classes')
```

Different classes

2.00 |

## ► Data preparation

```
[ ] ↳ 12 celdas ocultas
```

125 |

## ▼ 3. AI Model

```
1 from sklearn import model_selection
2 from sklearn.metrics import confusion_matrix
3
4 import tensorflow
5 import keras
6 from keras.layers import Dense, Dropout, Activation, Flatten, Conv1D, Conv2D, MaxPooling1D
7 from keras.utils import np_utils
8
9 from keras import models, layers, optimizers
10 from sklearn.model_selection import train_test_split
11 from sklearn.metrics import confusion_matrix, accuracy_score
12 from sklearn.utils import class_weight
13
14 from tensorflow.keras.optimizers import SGD, RMSprop, Adam, Adagrad, Adadelta, RMSprop
15 from keras.models import Sequential, model_from_json
16 from keras.preprocessing.image import ImageDataGenerator
17 from keras.callbacks import ReduceLROnPlateau, ModelCheckpoint
18 from keras import backend as K
19 from keras.applications.vgg16 import VGG16
20 from keras.models import Model
21
22 from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, la
23 import itertools
```

Define:

- Input length.
- Neurons in the last layer
- Batch\_size for training with SGD.

```
1 signal_length = 187
2 batch_size = 256
3 n_classes = 5
```

## ▼ Model

```

1 model = Sequential()
2
3 model.add(Conv1D(32, kernel_size=(5), input_shape=(signal_length, 1)))
4 model.add(Dropout(0.5))
5 model.add(BatchNormalization())
6 model.add(Activation('relu'))
7
8 model.add(Conv1D(32, (4)))
9 model.add(Dropout(0.5))
10 model.add(BatchNormalization())
11 model.add(Activation('relu'))
12 model.add(MaxPooling1D(pool_size=(2)))
13
14 model.add(Conv1D(32, (4)))
15 model.add(Dropout(0.5))
16 model.add(BatchNormalization())
17 model.add(Activation('relu'))
18 model.add(MaxPooling1D(pool_size=(2)))
19
20 model.add(Conv1D(32, (4)))
21 model.add(Dropout(0.5))
22 model.add(BatchNormalization())
23 model.add(Activation('relu'))
24 model.add(MaxPooling1D(pool_size=(2)))
25
26 model.add(Flatten())
27
28 model.add(Dense(128, activation='sigmoid'))
29 model.add(Dropout(0.5))
30 model.add(Dense(n_classes, activation='softmax'))
31
32 model.summary()
33
34 model.compile(loss=keras.losses.categorical_crossentropy,
35               optimizer=tensorflow.keras.optimizers.Adadelta(),
36               metrics=['accuracy'])
37

```

conv1d (Conv1D)	(None, 183, 32)	192
dropout (Dropout)	(None, 183, 32)	0
batch_normalization (Batch Normalization)	(None, 183, 32)	128
activation (Activation)	(None, 183, 32)	0
conv1d_1 (Conv1D)	(None, 180, 32)	4128
dropout_1 (Dropout)	(None, 180, 32)	0
batch_normalization_1 (Batch Normalization)	(None, 180, 32)	128
activation_1 (Activation)	(None, 180, 32)	0
max_pooling1d (MaxPooling1D)	(None, 90, 32)	0

max_pooling1d (MaxPooling1D)	(None, 90, 32)	0
conv1d_2 (Conv1D)	(None, 87, 32)	4128
dropout_2 (Dropout)	(None, 87, 32)	0
batch_normalization_2 (Batch Normalization)	(None, 87, 32)	128
activation_2 (Activation)	(None, 87, 32)	0
max_pooling1d_1 (MaxPooling1D)	(None, 43, 32)	0
conv1d_3 (Conv1D)	(None, 40, 32)	4128
dropout_3 (Dropout)	(None, 40, 32)	0
batch_normalization_3 (Batch Normalization)	(None, 40, 32)	128
activation_3 (Activation)	(None, 40, 32)	0
max_pooling1d_2 (MaxPooling1D)	(None, 20, 32)	0
flatten (Flatten)	(None, 640)	0
dense (Dense)	(None, 128)	82048
dropout_4 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 5)	645

=====

Total params: 95,781  
Trainable params: 95,525  
Non-trainable params: 256

To compile the model, `.compile()` is called. Here we specify which loss function we use, which optimizer and which metrics we want to keep for each epoch.

```
1 model.compile(loss='categorical_crossentropy', optimizer=Adam(), metrics=['accuracy'])
```

We train the model a number of epochs and with a specified `batch_size`. This returns us a history object with the accuracy of all the training phases.

## ▼ Training

```
1 history = model.fit(X_train, y_train,
2                     epochs=75,
3                     batch_size=batch_size,
```

```
3         batch_size=batch_size,  
4         verbose=1,  
5         validation_data=(X_test, y_test))
```

```
2/4/2/4 [=====] - 7s 25ms/step - loss: 0.0968 - accuracy:  
Epoch 47/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0965 - accuracy:  
Epoch 48/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0948 - accuracy:  
Epoch 49/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0951 - accuracy:  
Epoch 50/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0969 - accuracy:  
Epoch 51/75  
274/274 [=====] - 7s 24ms/step - loss: 0.0950 - accuracy:  
Epoch 52/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0958 - accuracy:  
Epoch 53/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0953 - accuracy:  
Epoch 54/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0937 - accuracy:  
Epoch 55/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0928 - accuracy:  
Epoch 56/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0935 - accuracy:  
Epoch 57/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0956 - accuracy:  
Epoch 58/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0901 - accuracy:  
Epoch 59/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0919 - accuracy:  
Epoch 60/75  
274/274 [=====] - 7s 24ms/step - loss: 0.0930 - accuracy:  
Epoch 61/75  
274/274 [=====] - 7s 24ms/step - loss: 0.0925 - accuracy:  
Epoch 62/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0887 - accuracy:  
Epoch 63/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0897 - accuracy:  
Epoch 64/75  
274/274 [=====] - 7s 26ms/step - loss: 0.0901 - accuracy:  
Epoch 65/75  
274/274 [=====] - 7s 24ms/step - loss: 0.0900 - accuracy:  
Epoch 66/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0885 - accuracy:  
Epoch 67/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0909 - accuracy:  
Epoch 68/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0889 - accuracy:  
Epoch 69/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0900 - accuracy:  
Epoch 70/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0888 - accuracy:  
Epoch 71/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0883 - accuracy:  
Epoch 72/75  
274/274 [=====] - 7s 25ms/step - loss: 0.0875 - accuracy:  
Epoch 73/75  
274/274 [=====] - 7s 24ms/step - loss: 0.0882 - accuracy:  
Epoch 74/75  
274/274 [=====] - 7s 24ms/step - loss: 0.0857 - accuracy:
```



## Access to the historical accuracy of the model (with the history attribute)

1 history.history

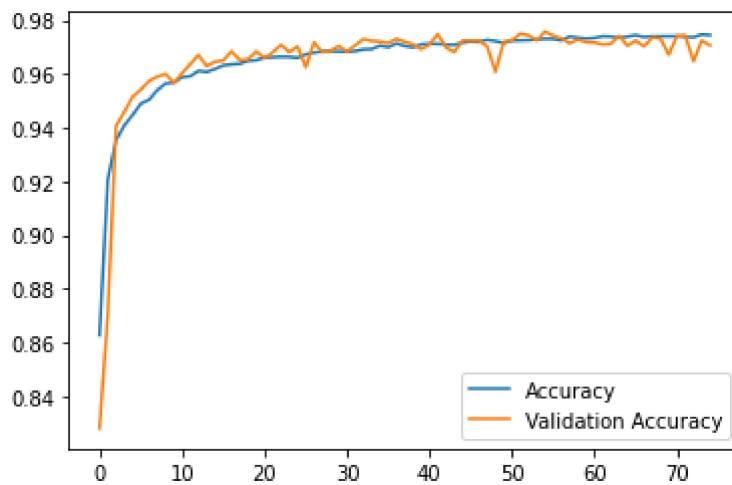
0.11635860055685043,  
 0.12120477110147476,  
 0.11769509315490723,  
 0.11101412028074265,  
 0.11513130366802216,  
 0.10697878152132034,  
 0.10440034419298172,  
 0.10927128791809082,  
 0.10051724314689636,  
 0.12406480312347412,  
 0.09744863212108612,  
 0.10840492695569992,  
 0.10468432307243347,  
 0.10148876160383224,  
 0.10558591037988663,  
 0.09736981987953186,  
 0.09891257435083389,  
 0.09821584075689316,  
 0.09725581854581833,  
 0.10032787919044495,  
 0.09103266149759293,  
 0.0982784628868103,  
 0.09699684381484985,  
 0.10702674835920334,  
 0.09707674384117126,  
 0.08697402477264404,  
 0.10201660543680191,  
 0.10459234565496445,  
 0.09405051916837692,  
 0.09666819125413895,  
 0.091598279774189,  
 0.09911223500967026,  
 0.14012511074543,  
 0.09638926386833191,  
 0.09862349182367325,  
 0.0911770761013031,  
 0.08987992256879807,  
 0.09612362831830978,  
 0.08746001869440079,  
 0.09349197149276733,  
 0.09117641299962997,  
 0.10377916693687439,  
 0.09539712220430374,  
 0.10085691511631012,  
 0.09267361462116241,  
 0.09856083989143372,  
 0.09981495141983032,  
 0.09145855158567429,  
 0.10609102994203568,  
 0.09644510596990585,  
 0.10398600995540619,  
 0.09261532127857208,  
 0.09721055051088333

```
0.09731055051088333,  
0.11862365156412125,  
0.08503658324480057,  
0.09194368124008179,  
0.12440980225801468,  
0.10360081493854523,  
0.1055551313001113011
```

To see if our model is overfitting, a graph is drawn with the accuracy in train and in validation using the data from the history object.

```
1 plt.plot(history.history['accuracy'], label='Accuracy')  
2 plt.plot(history.history['val_accuracy'], label='Validation Accuracy')  
3 plt.legend()
```


<matplotlib.legend.Legend at 0x7f58f19ee810>



And the Loss

```
1 plt.plot(history.history['loss'], label='Loss')  
2 plt.plot(history.history['val_loss'], label='Validation Loss' )  
3 plt.legend()
```

<matplotlib.legend.Legend at 0x7f58f31f7390>



## ▼ Model Test


```
1 y_pred = model.predict(X_test, batch_size=1000)
2 y_pred.shape
```

(17511, 5)

```
1 print(y_test[299])
2 plt.bar(range(n_classes), height=y_pred[4])
```

[1. 0. 0. 0. 0.]

<BarContainer object of 5 artists>



This code snippet generates a model report, and the next a confusion matrix.

```
1 print(classification_report(y_test.argmax(axis=1), y_pred.argmax(axis=1)))
```

	precision	recall	f1-score	support
0	0.97	1.00	0.98	14500
1	0.96	0.64	0.77	441
2	0.97	0.84	0.90	1137
3	0.81	0.34	0.48	124
4	1.00	0.94	0.97	1309
accuracy			0.97	17511
macro avg	0.94	0.75	0.82	17511
weighted avg	0.97	0.97	0.97	17511

```
1 def plot_confusion_matrix(cm, classes,
2                             normalize=False,
3                             title='Confusion matrix',
4                             cmap=plt.cm.Blues):
5     """
6     This function prints and plots the confusion matrix.
7     Normalization can be applied by setting `normalize=True`.
8     """
9     if normalize:
10         cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
```

```
11     print("Normalized confusion matrix")
12     else:
13         print('Confusion matrix, without normalization')
14
15     plt.imshow(cm, interpolation='nearest', cmap=cmap)
16     plt.title(title)
17     plt.colorbar()
18     tick_marks = np.arange(len(classes))
19     plt.xticks(tick_marks, classes, rotation=45)
20     plt.yticks(tick_marks, classes)
21
22     fmt = '.2f' if normalize else 'd'
23     thresh = cm.max() / 2.
24     for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
25         plt.text(j, i, format(cm[i, j], fmt),
26                 horizontalalignment="center",
27                 color="white" if cm[i, j] > thresh else "black")
28
29     plt.tight_layout()
30     plt.ylabel('True label')
31     plt.xlabel('Predicted label')
32
33 # Compute confusion matrix
34 cnf_matrix = confusion_matrix(y_test.argmax(axis=1), y_pred.argmax(axis=1))
35 np.set_printoptions(precision=2)
36
37 # Plot non-normalized confusion matrix
38 plt.figure(figsize=(10, 10))
39 plot_confusion_matrix(cnf_matrix, classes=['N', 'S', 'V', 'F', 'Q'],
40                       title='Confusion matrix, with normalization',
41                       normalize=True)
42 plt.show()
```

Normalized confusion matrix

